SQL Bootcamp

Wifi GA_GUEST

Pass: yellowpencil

Head to:

Github.com/morrisdata

Download

SQL bootcamp repository

Follow directions to install PGADMIN4 *note work devices may not work



Value of Data

Objectives

- Set up Environment
- Pgadmin4
- Postgres DB
- Connect to AWS
- Connect to Local Database
- Data Flow & SQL
- Retrieving and filtering data with SQL
- Aggregations with SQL
- Dynamic data referencing with SQL (joins)
- Demo 3 Value Logic/Case Statements/Subselects
- Next steps

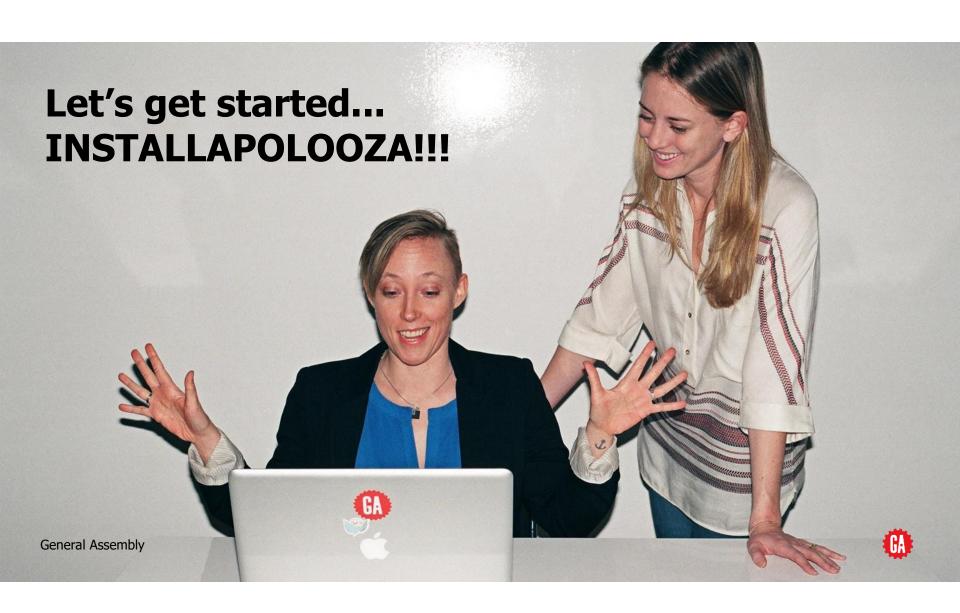
GENERAL ASSEMBLY





Matthew Morris

Advanced Analytics Data Analyst, Costco Data Analytics Course Instructor, General Assembly



nstallations/Config

- Load PGADMIN and connect to AWS
- Load Postgres(bonus material not needed for 99% of todays class)
- Help fellow students connect if you finish
- As a class lets load a table together
- If you cannot get connected or configured you can group up with another stude.
 Or reschedule for another class time. Installs and configuration troubleshooting
 Can depend on many factors and we wont have time to troubleshoot everyones
 Devices



CREATING A LOCAL DATABASE

POSTGRES AND PGADMIN INSTALL

BRIEF OVERVIEW DATA MODELING



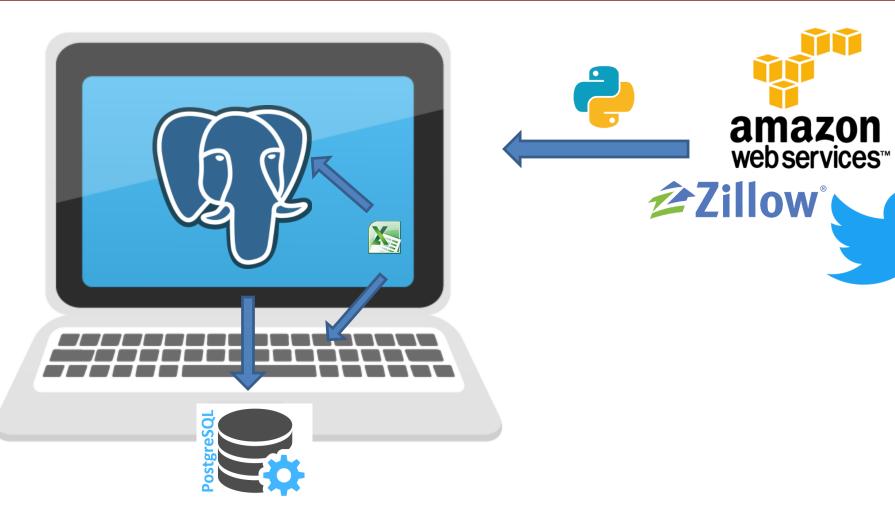
Matthew Morris

Git: Morrisdata

 ${\it Matthew Morris.DA@gmail.com}$

CREATING LOCAL DATABASES

OBTAINING AND MODELING DATA

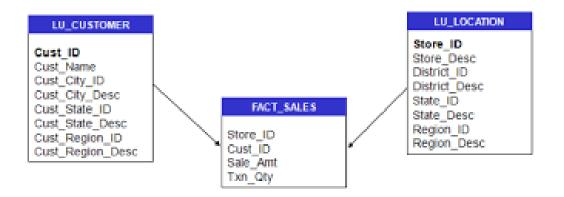


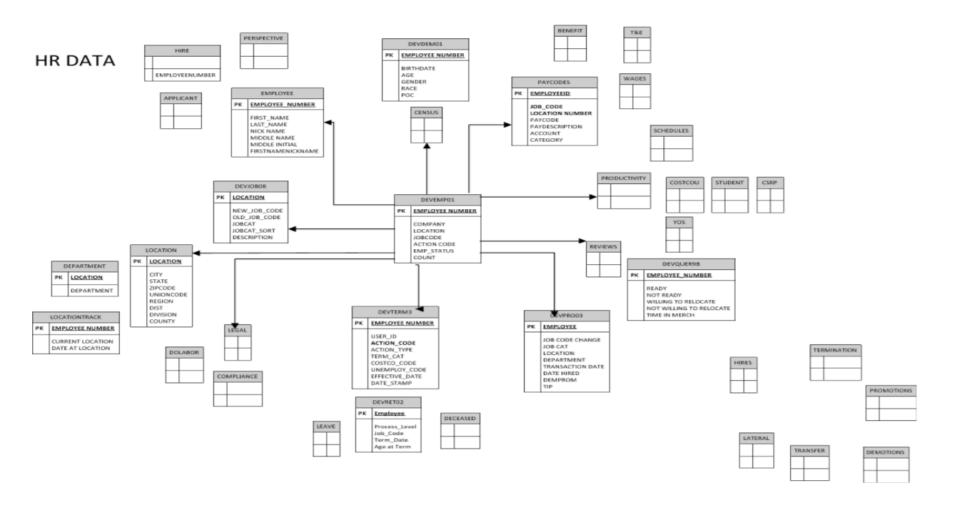


Id	Release date	Record Label	Artist	Song	Album	sales
1						
2						
3						
4						
5						

Id	Artist	Song	Album	sales
1				
2				
3				
4				
5				

SALES		
FIELD ID ARTIST SONG ALBUN	Char	LENGTH 1 25 225 225





Creating Table and loading Data

```
CREATE TABLE
schema.tablename
    field1
            integer,
            numeric,
    field2
            character(30),
    field3
   field4
            money
```

Creating Table and Loading Data

COPY table FROM 'datasource' Delimiter ',' filetype Header;



FUNDAMENTALS OF DATAFLOW AND SQL

Explain where SQL fits in the dataflow Retrieve and filter data with basic SQL Navigate a Relational Database

SELECT

FROM

WHERE

ORDERBY

LIMIT

Matthew Morris

Git: Morrisdata

Matthew Morris. DA@gmail.com



FUNDAMENTALS OF DATAFLOW AND SQL

GATHERING REQUIREMENTS

GATHERING REQUIREMENTS



FUNDAMENTALS OF DATAFLOW AND SQL

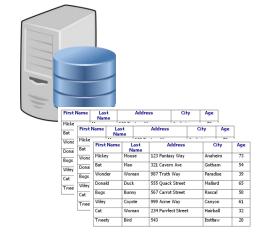
GATHERING REQUIREMENTS



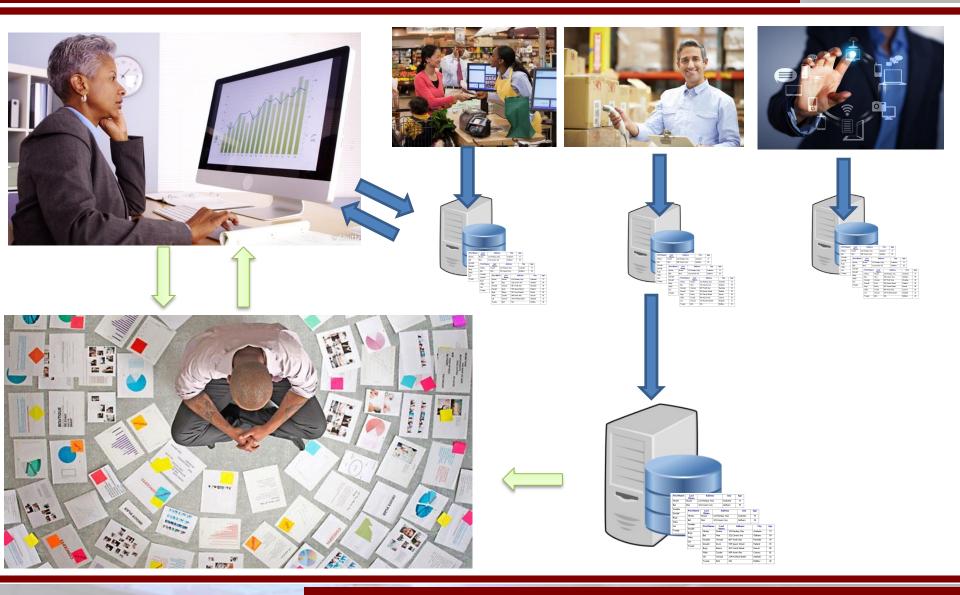


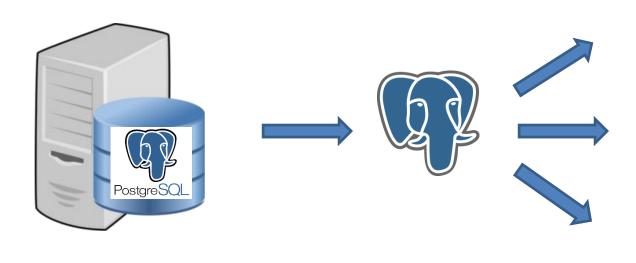






GATHERING REQUIREMENTS

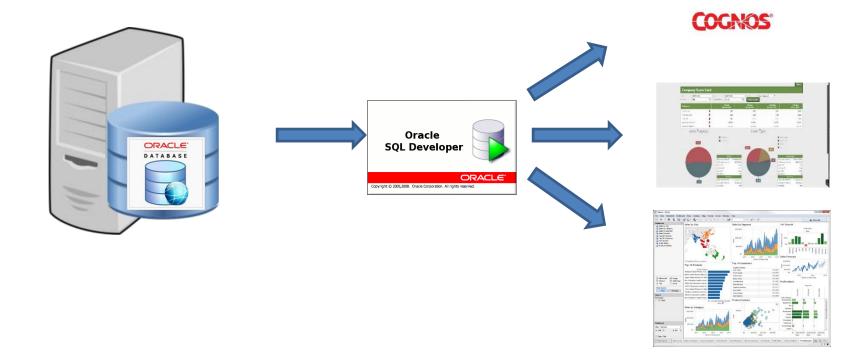


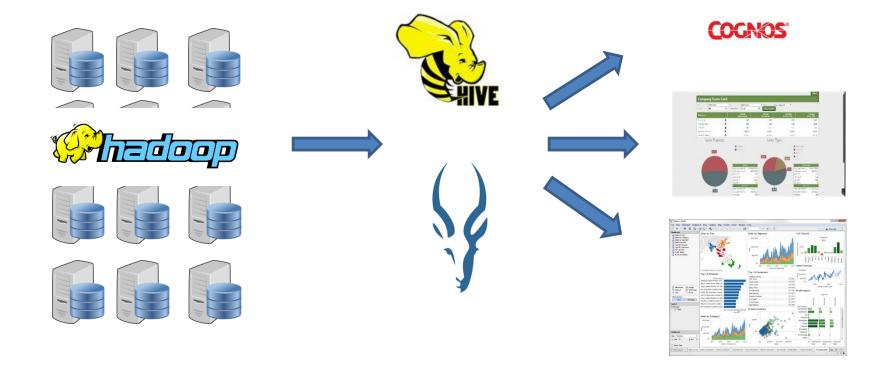


COGNOS











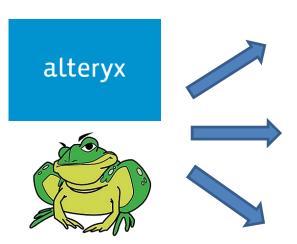










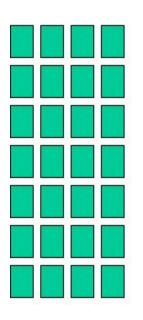








Relational DB



Libraries/Collections

Schema

Tables/Files/Objects

Members/Partitions in files



Review

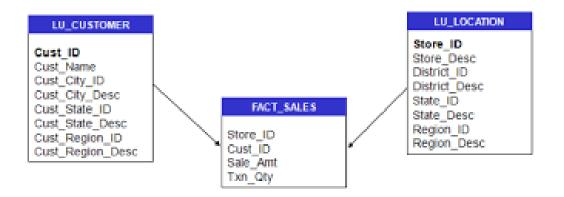


Id	Release date	Record Label	Artist	Song	Album	sales
1						
2						
3						
4						
5						

Id	Artist	Song	Album	sales
1				
2				
3				
4				
5				

Review

SALES		
FIELD	TYPE	LENGTH
ID	PK	1
ARTIST	Char	25
SONG	Char	225
ALBUM	Char	225



Relational DB

NETWORK

Contacts

Name

Friends

Family

Business Name

Address

Phone number

Occupation

Gift Ideas

Name

Gift ideas

Gifts received

Gifts given

Party

Name

Holidays

Birthday

Anniversaries



<u>Birthdays</u>

Name

Birthday

Gift Ideas

Gifts received

Gifts given

Family

Relational DB



© geology.com

Understanding your data. What are the tables? What are the fields?

How might you link the tables?



SELECT FROM WHERE ORDER BY LIMIT

Commenting



-- Basic commenting

/* Multiple line comment

End of Multiple line comment*/

SELECT *

SELECT FIELD1, FIELD2 ...

SELECT (FIELD1+FIELD2), FIELD 3...

SELECT SUM(FIELD1), FIELD2

CAST(field AS type)

A helpful String function

Is Case Cost a discount for stores or for customers? Use SQL and field exploration and math in your statement to answer this question.

SELECT DISTINCT Location, NumberOfSales

Location NumberOfSales

Seattle 101

Seattle 40

Tacoma 72

SELECT DISTINCT Location, NumberOfSales, Date

Location	NumberOfSales	Date
Seattle	101	10/28/17
Seattle	101	10/27/17
Seattle	40	10/26/17
Tacoma	72	10/28/17
Tacoma	72	10/27/17

WHERE COUNTRY = US

WHERE COUNTRY = US AND STATE = WA

WHERE COUNTRY = US AND STATE = WA OR SALES > 100

ORDER BY 1

ORDER BY 1,2 DESC

LIMIT 1000

ROWNUM <= 1000

- 1. Select various fields from the SALES table that interest you.
 - *BE sure to use LIMIT 1000
- 2. Practice using filters.
- 3. Use AND to apply multiple filters Change the sort.
- 4. Save your Query

Use your new skills to review Iowa Liquor Sales





Filters and Aggregations

FILTERS = , !=, >, <
IN, NOT IN, BETWEEN, LIKE, NOT LIKE
SUM, MIN, MAX, COUNT
GROUP BY, HAVING
COMMENTING



Git: Morrisdata MatthewMorris.DA@gmail.com





SELECT FROM WHERE **GROUP BY HAVING ORDER BY LIMIT**

WHERE

- =, !=, >,<
- IS NULL, IS NOT NULL
- IN, NOT IN
- BETWEEN
- LIKE
- OR



Aggregate functions

- MIN
- MAX
- SUM
- COUNT



GROUP BY store, item

GROUP BY 1,2

HAVING AVG(sales)>100 AND COUNT(customers)>20

WHERE – filter for dimensions and measures

GROUP BY –groups dimensions when a measure is aggregated

HAVING – filter for aggregated measure

SELECT

- Fields you want to see in your results

FROM

- Table where fields come from

WHERE

- Filters for your results

GROUP BY

- Groups dimensions when using an aggregate

HAVING

- Filters aggregations

ORDER BY

- How you can sort your results

LIMIT

- Limits number of records returned

Fix the Code

```
SELECT Store, (cost –sell price), SUM(sales),
FROM sales
WHERE Category = 'Tequila'
AND units purchased >2
GROUP BY Store
HAVING SUM(sales) > 30.00
ORDER BY 3
```

Workshop

Which products have a case cost of more than \$100?

Which tequilas have a case cost of more than \$100?

Which tequilas or scotch whiskies have a case cost of more than \$100?

Which tequilas or scotch whiskies have a cast cost between \$100 and \$120?

Which whiskies of any kind cost more than \$100?'

Which whiskies of any kind cost between \$100 and \$150?

Which products except tequilas cost between \$100 and \$120?



Querying Relational Database

UNION
JOIN 2 Tables
JOIN Multiple Tables

Matthew Morris

Git: Morrisdata
MatthewMorris.DA@gmail.com



Query a Relational Database



SELECT FROM

JOIN

ON

WHERE

GROUP BY

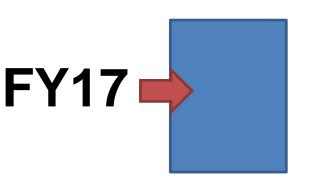
HAVING

UNION

ORDER BY

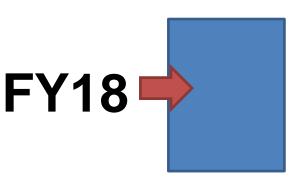
LIMIT

Unions



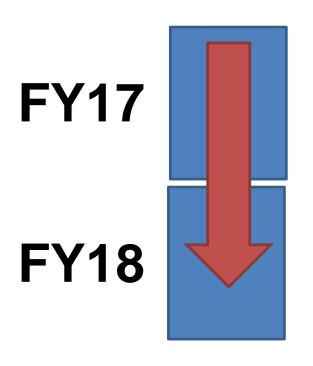
SELECT fy, pd, store_name, week1, week2, week3 week4 FROM FY17

UNION



SELECT fy, pd, store_name, week1, week2, week3 week4 FROM FY18

Unions

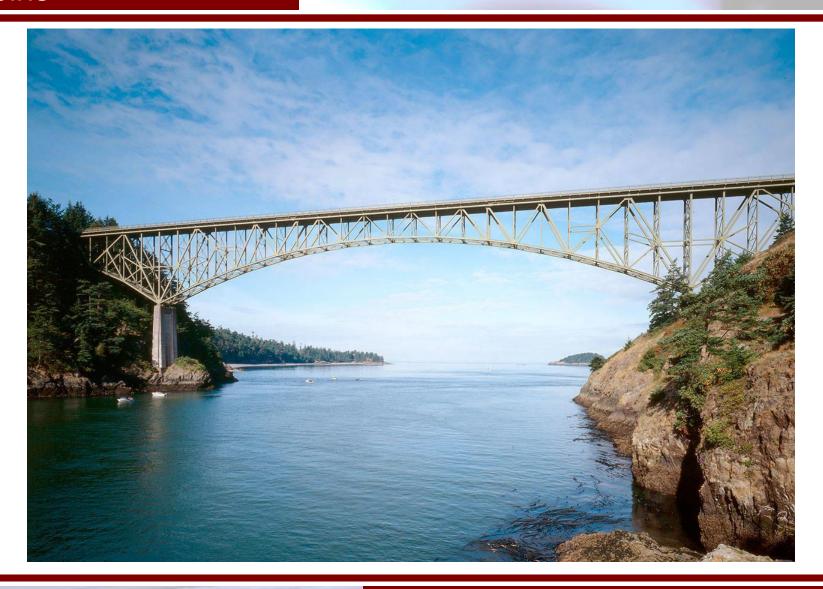


SELECT fy, pd, store_name, week1, week2, week3 week4
FROM FY17
UNION
SELECT fy, pd, store_name, week1, week2, week3 week4
FROM FY18

COLUMNS CONDITIONS UNION and UNION ALL ORDER BY

QUERY A RELATIONAL DATABASE

JOIN 1 Table



Querying a Relational Database



Query a Relational Database

LEFT/PRIMARY

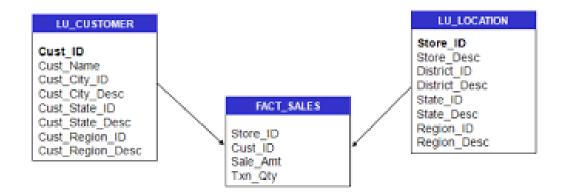


What table is the transaction table?

If you wanted to link on the lowest level of detail to the other tables what fields would you use?

SALES		
FIELD	TYPE	LENGTH
ID	PK	1
ARTIST	Char	25
SONG	Char	225
ALBUM	Char	225

Create a rough sketch with how The Iowa Liquor Sales Database Would JOIN



SELECT a.item, b.description, a.sales

FROM sales a

JOIN products b

ON a.item=b.item

- Create separate queries to join each table to Sales
 - a. Products to Sales
 - b. County to Sales
 - c. Stores to Sales
- 2. Use this as an opportunity to bring fields in from both tables.
- Try out some aggregations or Wild card searches. Stretch with an Aggregate and a Group by

SELECT c.field, a.field, b.field, a.field, c.field

FROM table 1 a

JOIN table 2 b

ON a.field=b.field

JOIN table3 c

ON a.field=c.field



Join types

LEFT OUTER
RIGHT OUTER
LEFT EXCEPTION
RIGHT EXCEPTION
CROSS
COALESCE



Matthew Morris

Git: Morrisdata
MatthewMorris.DA@gmail.com

LEFT/PRIMARY



Types of Joins

nner Join

.eft-Outer Join

Right-Outer Join

Exception Join

Right-Exception Join

Cross Join

Match in both tables

Includes data from the primary table that may not have matches

Includes data from the secondary table that may not have matches

Returns Primary table data that does not match with the secondary table

Returns Secondary table data that does not match with the Primary table

Returns all data whether a match exists or not

OUTER, EXCEPTION, AND CARTESIAN JOINS

Employees

id	first_name	last_name
2	Gabe	Moore
3	Doreen	Mandeville
5	Simone	MacDonald
7	Madisen	Flateman
11	Ian	Paasche
13	Mimi	St. Felix

Salaries

id	current_salary
2	50000
3	60000
7	55000
11	75000
13	120000
17	70000

OUTER, EXCEPTION, AND CARTESIAN JOINS

- An INNER JOIN (also called a direct join) displays only the rows that have a match in both joined tables.
- An INNER JOIN would yield this table:

SELECT * FROM employees INNER JOIN salaries ON employees.ID = salaries.ID;

id	first_name	last_name	id	current_salary
2	Gabe	Moore	2	50000
3	Doreen	Mandeville	3	60000
7	Madisen	Flateman	7	55000
11	Ian	Paasche	11	75000
13	Mimi	St. Felix	13	7000

OUTER, EXCEPTION, AND CARTESIAN JOINS

A LEFT OUTER JOIN returns both:

- Data that both tables have in common.
- Data from the **primary** table selected, which does
 not have matching data to
 join to in the secondary
 table.
- A LEFT OUTER JOIN would yield this table:

SELECT * FROM employees LEFT OUTER JOIN salaries ON employees.ID = salaries.ID;

id	first_name	last_name	id	current_salary
2	Gabe	Moore	2	50000
3	Doreen	Mandeville	3	60000
5	Simone	MacDonald	NULL	NULL
7	Madisen	Flateman	7	55000
11	Ian	Paasche	11	75000
13	Mimi	St. Felix	13	120000

OUTER, EXCEPTION, AND CARTESIAN JOINS

- A RIGHT OUTER JOIN returns both:
 - Data that two tables have in common.
 - Data from the **secondary** table selected, which does
 not have matching data to
 join to in the primary table.
- A RIGHT OUTER JOIN would yield this table:

SELECT * FROM employees RIGHT OUTER
JOIN salaries ON employees.ID = salaries.ID;

id	first_name	last_name	id	current_salary
2	Gabe	Moore	2	50000
3	Doreen	Mandeville	3	60000
7	Madisen	Flateman	7	55000
11	Ian	Paasche	11	75000
13	Mimi	St. Felix	13	120000
NULL	NULL	NULL	17	70000

OUTER, EXCEPTION, AND CARTESIAN JOINS

- A FULL OUTER JOIN returns all data from each table, regardless of whether it has matching data in the other table.
- A FULL OUTER JOIN would yield this table:

SELECT * FROM employees FULL OUTER

JOIN salaries ON employees.ID = salaries.ID;

id	first_name	last_name	id	current_salary
2	Gabe	Moore	2	50000
3	Doreen	Mandeville	3	60000
5	Simone	MacDonald	NULL	NULL
7	Madisen	Flateman	7	55000
11	Ian	Paasche	11	75000
13	Mimi	St. Felix	13	120000
NULL	NULL	NULL	17	70000

OUTER, EXCEPTION, AND CARTESIAN JOINS

- An EXCEPTION JOIN
 returns only the data from
 the primary, or first table
 selected, which does not
 have matching data to join
 to in the secondary table.
- An EXCEPTION JOIN would yield this table:

SELECT * FROM employees LEFT OUTER

JOIN salaries ON employees.ID = salaries.ID

WHERE salaries.ID IS NULL;

id	first_name	last_name	id	current_salary
5	Simone	MacDonald	NULL	NULL

OUTER, EXCEPTION, AND CARTESIAN JOINS

- A RIGHT EXCEPTION

 JOIN returns only the data from the secondary, or second table selected, which does not have matching data to join to in the primary table.
- A RIGHT EXCEPTION
 JOIN would yield this table:

SELECT * FROM employees RIGHT OUTER
JOIN salaries ON employees.ID = salaries.ID
WHERE employees.ID IS NULL;

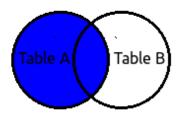
id	first_name	last_name	id	current_salary
NULL	NULL	NULL	17	70000

OUTER, EXCEPTION, AND CARTESIAN JOINS

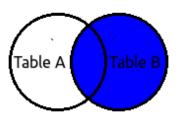
- A CROSS JOIN matches every row of the primary table with every row of the secondary table.
- This type of join results in a
 Cartesian product of the tables,
 is generally detrimental to slow
 performance, and is not desired.
- A CROSS JOIN would yield this table:

SELECT * FROM employees CROSS JOIN salaries;

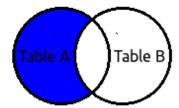
id	first_name	last_name	id	current_salary
2	Gabe	Moore	2	50000
3	Doreen	Mandeville	2	50000
5	Simone	MacDonald	2	50000
7	Madisen	Flateman	2	50000
11	lan	Paasche	2	50000
13	Mimi	St. Felix	2	50000
2	Gabe	Moore	3	60000
3	Doreen	Mandeville	3	60000
5	Simone	MacDonald	3	60000
7	Madisen	Flateman	3	60000
11	lan	Paasche	3	60000
13	Mimi	St. Felix	3	60000
2	Gabe	Moore	7	55000
3	Doreen	Mandeville	7	55000
5	Simone	MacDonald	7	55000
7	Madisen	Flateman	7	55000
11	lan	Paasche	7	55000
13	Mimi	St. Felix	7	55000
2	Gabe	Moore	11	75000
3	Doreen	Mandeville	11	75000
5	Simone	MacDonald	11	75000
7	Madisen	Flateman	11	75000
11	lan	Paasche	11	75000
13	Mimi	St. Felix	11	75000
2	Gabe	Moore	13	120000
3	Doreen	Mandeville	13	120000
5	Simone	MacDonald	13	120000
7	Madisen	Flateman	13	120000
11	lan	Paasche	13	120000
13	Mimi	St. Felix	13	120000
2	Gabe	Moore	17	70000
3	Doreen	Mandeville	17	70000
5	Simone	MacDonald	17	70000
7	Madisen	Flateman	17	70000
11	lan	Paasche	17	70000
13	Mimi	St. Felix	17	70000



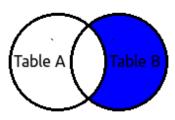
SELECT [list] FROM
[Table A] A
LEFT JOIN
[Table B] B
ON A.Value = B.Value



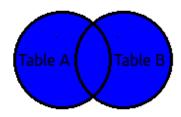
SELECT [list] FROM
[Table A] A
RIGHT JOIN
[Table B] B
ON A.Value = B.Value



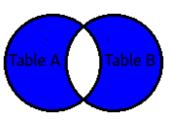
SELECT [list] FROM
[Table A] A
LEFT JOIN
[Table B] B
ON A.Value = B.Value
WHERE B.Value IS NULL



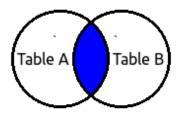
SELECT [list] FROM
[Table A] A
RIGHT JOIN
[Table B] B
ON A.Value = B.Value
WHERE A.Value IS NULL



SELECT [list] FROM
[Table A] A
FULL OUTER JOIN
[Table B] B
ON A.Value = B.Value



SELECT [list] FROM
[Table A] A
FULL OUTER JOIN
[Table B] B
ON A.Value = B.Value
WHERE A.Value IS NULL
OR B.Value IS NULL



SELECT [list] FROM
[Table A] A
INNER JOIN
[Table B] B
ON A.Value = B.Value

"I want to see all of the information we can get on inactive stores for sales, if there are any, and their addresses."

Process to pick the right join (be methodical)



Common Table Expressions/Subselect

WITH SELECT INTO



Matthew Morris

Git: Morrisdata MatthewMorris.DA@gmail.com



Previously in Data Analytics

LEFT OUTER
RIGHT OUTER
LEFT EXCEPTION
RIGHT EXCEPTION
CROSS
COALESCE





WITH

SELECT INTO

FROM

JOIN

ON

WHERE

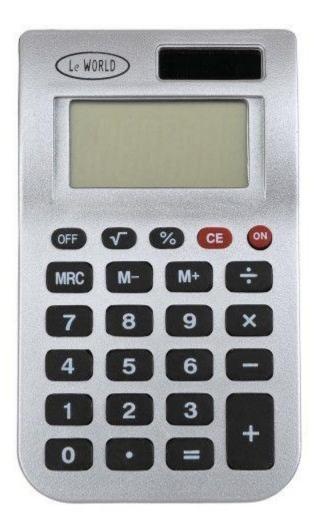
GROUP BY

HAVING

UNION

ORDER BY

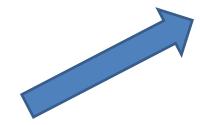
LIMIT



A common table expression (CTE) can be thought of as a temporary result set that is defined within the execution scope of a single SELECT, INSERT, UPDATE, DELETE, or CREATE VIEW statement. A CTE is similar to a derived table in that it is not stored as an object and lasts only for the duration of the query. Unlike a derived table, a CTE can be self-referencing and can be referenced multiple times in the same query.

SELECT county, COUNT(DISTINCT store) AS NumOfStores, SUM(total) AS total_sales FROM sales
WHERE category_name LIKE '%TEQUILA%'
GROUP BY county, store
HAVING count(store)>2 AND sum(total)<10000

	1		
county	numofstores	¥	total_sales 💌
Adair		1	7659.95
Adair		1	2420.24
Adair		1	2123.88
Adair		1	1297.44
Adair		1	2165.28
Adair		1	308.19
Adams		1	5710.36
Allamakee		1	9938.2
Allamakee		1	8080.96
Allamakee		1	1047.16



SELECT county, SUM(NumOfStores)
FROM county_sales
GROUP BY county



GROUP BY county

```
WITH county_sales AS (

SELECT county, COUNT(DISTINCT store) AS NumOfStores, SUM(total) AS total_sales
FROM sales
WHERE category_name LIKE '%TEQUILA%'
GROUP BY county, store
HAVING count(store)>2 AND sum(total)<10000
)

SELECT county,
SUM(NumOfStores)
FROM county sales

County ■ Num_of_Stores ■
```

County -	Num_of	_Stores	¥
Wayne			2
[null]			8
Poweshiel			9
Guthrie			4
Delaware			3
Harrison			4
Fayette			6
Winnebag			4
Sioux			5
Greene			7

```
SELECT county, SUM(NumOfStores) FROM
```

(SELECT county, store, COUNT(DISTINCT store) AS NumOfStores, SUM(total) AS total_sales

FROM sales

WHERE Category_name LIKE '%TEQUILA%'

GROUP BY county, store

HAVING count(store)>20 AND sum(total)>500000) AS county_sales

GROUP BY county

County 💌	Num_of_Stores	¥
Wayne		2
[null]		8
Poweshiel		9
Guthrie		4
Delaware		3
Harrison		4
Fayette		6
Winnebag		4
Sioux		5
Greene		7

Common Table Expressions

DEMO(aggregate over an aggregate)

ACTIVITY: AGGREGATE OVER AN AGGREGATE



What is the total sum of stores by county?

WITH a AS (SELECT county, population FROM counties WHERE population >5000),

b AS (SELECT county, sum(total) as total_sales FROM sales GROUP BY county)

a

county	population 💌
Adair	7682
Allamakee	14330
Appanoose	12884
Audubon	6119
Benton	26076
Black Hawk	131090
Boone	26306
Bremer	24276
Buchanan	20958
Buena Vista	20260

b

county	total_sales 💌
Wayne	98824.4
Montgomery	839109.2
	193443.2
Poweshiek	2012519.27
Story	12267027.18
Calhoun	482178.34
Dallas	7898228.42
Decatur	248492.15
Worth	339968.23
Carroll	3183557.2

SELECT *
FROM a
JOIN b
ON a.county = b.county

population 💌	county2	total_sales 💌
7682	Adair	603088.67
14330	Allamakee	1107480.66
12884	Appanoose	1089064.58
6119	Audubon	241821.88
26076	Benton	994344.98
131090	Black Hawk	22967283.29
26306	Boone	2441399.88
24276	Bremer	2507674.79
20958	Buchanan	1949162.4
20260	Buena Vista	2432925.21
	7682 14330 12884 6119 26076 131090 26306 24276 20958	7682 Adair 14330 Allamakee 12884 Appanoose 6119 Audubon 26076 Benton 131090 Black Hawk 26306 Boone 24276 Bremer 20958 Buchanan 20260 Buena Vista

WITH
a AS (SELECT county, population
FROM counties
WHERE population >5000),

b AS (SELECT county, sum(total) as total_sales FROM sales

GROUP BY county)

SELECT *
FROM a
JOIN b
ON a.county = b.county

county	population 🔻	county2	total_sales 💌
Adair	7682	Adair	603088.67
Allamakee	14330	Allamakee	1107480.66
Appanoose	12884	Appanoose	1089064.58
Audubon	6119	Audubon	241821.88
Benton	26076	Benton	994344.98
Black Hawk	131090	Black Hawk	22967283.29
Boone	26306	Boone	2441399.88
Bremer	24276	Bremer	2507674.79
Buchanan	20958	Buchanan	1949162.4
Buena Vista	20260	Buena Vista	2432925.21

SELECT * FROM

(SELECT county, population FROM counties WHERE population >5000) as a

JOIN

(SELECT county, sum(total) as total_sales FROM sales GROUP BY county) as b

ON a.county = b.county

county	population 💌	county2	total_sales 💌
Adair	7682	Adair	603088.67
Allamakee	14330	Allamakee	1107480.66
Appanoose	12884	Appanoose	1089064.58
Audubon	6119	Audubon	241821.88
Benton	26076	Benton	994344.98
Black Hawk	131090	Black Hawk	22967283.29
Boone	26306	Boone	2441399.88
Bremer	24276	Bremer	2507674.79
Buchanan	20958	Buchanan	1949162.4
Buena Vista	20260	Buena Vista	2432925.21

Common Table Expressions

DEMO(joining temp tables)

ACTIVITY: JOINING temp Tables



Create a process that looks at store and store count
LIMIT to 5 for a sample data set
We also want to review a sum of sales for anything like tequila
USE CTE's for this as we may want some more stats to add later



